

# Project #2 Big Data Tools

Gabriel de O. F. Gonçalves (uc2023238703)

João Pereira (uc2023243538)

Faculdade de Ciências e Tecnologia da Universidade de Coimbra

IACD - MECD

## 1 Kafka (Assigment 1)

- Before any task, start Zookeeper and Kafka:

```
sudo systemctl start zookeeper
sudo systemctl start kafka
```

### 1.1 Task1

- Create a kafka topic called task1-topic with one partition:

```
kafka-topics --create --topic task1-topic --bootstrap-server localhost:9092 --partitions
1 --replication-factor 1
```

#### 1.1.1 Kafka producer

- Initialize the Kafka producer:

```
producer = KafkaProducer( bootstrap_servers='localhost:9092',
value_serializer=lambda v: json.dumps(v).encode('utf-8'))
```

- Define the random data:

```
sensor_id = 1 temperature_reading = "sensor_id": sensor_id,
"temperature": round(random.uniform(15.0, 35.0), 1)
```

- Send Data to Kafka Topic:

```
producer.send('task1-topic', value=temperature_reading)
print(f"Sent: temperature_reading")
```

#### 1.1.2 Kafka consumer

- Set up Kafka consumer:

```
consumer = KafkaConsumer( 'task1-topic',
bootstrap_servers='localhost:9092',
value_deserializer=lambda x: json.loads(x.decode('utf-8')) )
```

- Consumer listens to the topic and continuously retrieves messages:

```
for message in consumer:
print(f"Received message: message.value")
```

### 1.2 Task 2

- Create a kafka topic called task2-topic with 3 partitions:

```
kafka-topics --create --topic task1-topic --bootstrap-server localhost:9092 --partitions
3 --replication-factor 1
```

### 1.2.1 Kafka producer

- **Initialize the Kafka producer:**  

```
producer = KafkaProducer( bootstrap_servers='localhost:9092',  
value_serializer=lambda v: json.dumps(v).encode('utf-8') )
```
- **Define users and activities:**  

```
user_ids = [f"user{i}" for i in range(1, 5)]  
activities = ["login", "logout", "purchase", "browse", "signup"]
```
- **Generate and send user activity logs:**  

```
activity_log = {  
    "user_id": random.choice(user_ids),  
    "activity": random.choice(activities) }  
producer.send('task2-topic', value=activity_log)  
print(f"Sent: activity_log")
```

### 1.2.2 Kafka consumer

- **Set up Kafka consumer to listen to the topic task2-topic:**  

```
consumer = KafkaConsumer( 'task2-topic',  
bootstrap_servers='localhost:9092',  
group_id=group_id,  
auto_offset_reset='earliest',  
value_deserializer=lambda v: json.loads(v.decode('utf-8')) )
```
- **Consume and Process Messages:**  

```
for message in consumer:  
    print(f"Consumer {consumer_id} received: {message.value}")
```

### 1.2.3 Multiple activated consumers behavior

- **Three consumers in activity-group are active:**  
Each consumer consumes from 1 partition
- **Two consumers in activity-group are active:**  
One consumer consumes from 2 partitions and the other consumes from the left partition
- **Four consumers in activity-group are active:**  
Only 3 consumers are able to consume and the 4th consumer doesn't consume from any partition

## 1.3 Task 3

- **Create Kafka Topics: Purchase and User-Activity:**  

```
'kafka-topics --create --topic purchase-topic --bootstrap-server localhost:9092 --partitions  
1 --replication-factor 1'  
'kafka-topics --create --topic user-activity-topic --bootstrap-server localhost:9092  
--partitions 1 --replication-factor 1'
```

### 1.3.1 Kafka producer

- **Initialize the Kafka producer:**  

```
producer = KafkaProducer( bootstrap_servers='localhost:9092',  
value_serializer=lambda v: json.dumps(v).encode('utf-8') )
```
- **Define user actions:**  

```
items = ["book", "laptop", "phone", "headphones"]  
activities = ["login", "page_view", "add_to_cart", "checkout"]
```

- **Send Data to Kafka Topics:**  

```
producer.send('purchase-topic', purchase_data)
producer.send('user-activity-topic', activity_data)
```

### 1.3.2 Kafka consumer purchase-group and activity-group

- Both use the same logic.
- **Set up Kafka consumer to listen to the topic:**  

```
consumer = KafkaConsumer( 'purchase-topic',
bootstrap_servers='localhost:9092',
group_id=group_id,
value_deserializer=lambda v: json.loads(v.decode('utf-8')) )
```
- **Dictionary to keep track info:**  

```
user_total_spent = defaultdict(float)
```
- **Consume and Process Messages:**  

```
for message in consumer:
    purchase = message.value
    user_id = purchase["user_id"]
    amount = purchase["amount"]

    user_total_spent[user_id] += amount
    print(f"User {user_id} Total Spent:{user_total_spent[user_id]}")
```

## 2 Hadoop (Assignment 2)

### 2.1 Task 1

- **Run task 1:**  

```
python task1.py fakefriends.csv
```
- **Define a MRStep with the mapper and reducer:**  
Mapper: Processes each line of input data and emits age as the key and the number of friends as the value.  
Reducer: Aggregates the number of friends by age and calculates the average.
- **def mapper\_get\_average(self, \_, line):**  
Splits each line of input into fields: `userId, userName, age, numFriends`.  
Yields the age as the key and the number of friends as the value.
- **def reducer\_get\_average(self, age, values):**  
In each age value, sum the number of friends and count the number of people that have that age.  
Divide those two values to calculate the average of friends at a given age.

### 2.2 Task 2

- **Run task 2:**  

```
python task2.py 1800.csv
```
- **Define a MRStep with the mapper and reducer:**  
Mapper: Emits the weather station ID as the key and its temperature as the value.  
Reducer: Calculates the minimum temperature for each weather station.

- `def mapper.get_min(self, _, line):`  
Splits each input line into columns.  
Extracts the weather station ID (first column) and temperature (fourth column).  
Yields the weather station ID as the key and the temperature as the value.
- `def reducer.get_min(self, weather_station, values):`  
Receives all temperature values for each weather station.  
Calculates the minimum temperature.  
Yields the weather station ID and its minimum temperature.

## 2.3 Task 3

- **Run Task 3:**  
`python task3.py Book`
- Define two MRSteps, one with the mapper and reducer, the other with the reducer\_sort
- `def mapper(self, _, line):`  
Processes each line of text.  
Extracts words using a regular expression.  
Emits each word with a count of 1.
- `def reducer(self, word, counts):`  
Receives all counts for each word.  
Sums the counts to calculate the total frequency.  
Emit total count and word as a tuple to prepare for sorting.
- `def reducer_sort(self, _, word_count_pairs):`  
Receives all (total\_count, word) pairs as a list.  
Sorts the list in descending order of frequency.  
Emits each word and its frequency.

## 2.4 Task 4

- **Run Task 4:**  
`python task4.py customer-orders.csv`
- Define two MRSteps, one with the mapper and reducer\_get\_sum, the other with the reducer\_sort
- `def mapper(self, _, line):`  
Reads each input line.  
Extracts customerId and amount.  
yield both values.
- `def reducer.get_sum(self, customer_id, values):`  
Receives all amounts for each customer.  
Calculates the total amount spent by summing the values.  
Emits a tuple (total.amount, customerId).
- `def reducer_sort(self, _, customer_amount_pairs):`  
Receives a list of (total\_amount, customerId) tuples  
Sorts the list in descending order of total.amount Emits total amount spent sorted by each customer.

## 2.5 Terminal Outputs

Task 1	Task 2	Task 3	Task 4
<pre>36 "246.60" 37 "249.33" 38 "193.53" 22 "206.43" 23 "246.30" 24 "233.80" 25 "197.45" 64 "281.33" 65 "298.20" 66 "276.44" 67 "214.62" 42 "303.50" 43 "230.57" 44 "282.17" 45 "309.54" 26 "242.06" 27 "228.12" 28 "209.10" 54 "278.08" 55 "295.54"</pre>	<pre>"GM000010962" 0 "ITE00100554" -148 "EZE00100082" -135</pre>	<pre>"there" 12 "more" 11 "very" 11 "when" 11 "down" 11 "them" 11 "like" 11 "fire" 10 "teddy" 10 "from" 10 "again" 10 "head" 9 "time" 9</pre>	<pre>"68" "6375.45" "73" "6206.20" "39" "6193.11" "54" "6065.39" "71" "5995.66" "2" "5994.59" "97" "5977.19" "46" "5963.11" "42" "5696.84" "59" "5642.89" "41" "5637.62" "0" "5524.95" "8" "5517.24" "85" "5503.43" "61" "5497.48"</pre>

Figure 1: Outputs for all tasks in the assignment 2.

## 3 Spark RDDs (Assignment 3)

### 3.1 Task 1

- Configure Spark application.
- Load dataset.
- Split each line by comma.
- Transform to key-value pairs.
- Find the minimum temperature for each category.
- Collect and print results.

### 3.2 Tasks 2 and 3

- Configure Spark application.
- Load the book text file as an RDD.
- Define the line number to start processing (zero-indexed).
- Filter RDD to keep only lines starting from the specified line.
- Split lines into words, convert to lowercase, and count word frequency.
- Sort word counts by frequency in ascending order.
- Collect and print results.

### 3.3 Tasks 4 and 5

- Configure Spark application.
- Load the dataset (customer orders) as an RDD.
- Split each line by comma.
- Transform into key-value pairs (Customer\_ID, Amount).

- Compute the total amount spent by each customer by summing values.
- Sort the results by total amount in descending order.
- Collect and print results.

### 3.4 Tasks 6 and 7

- Configure Spark application.
- Load the graph dataset as an RDD.
- Count the frequency of each superhero ID.
- Load the names dataset as an RDD.
- Map names dataset to (ID, Name) pairs.
- Join the two datasets on superhero ID.
- Map to (Name, Frequency) and sort by frequency in ascending order.
- Collect and display the results.

### 3.5 Terminal Outputs

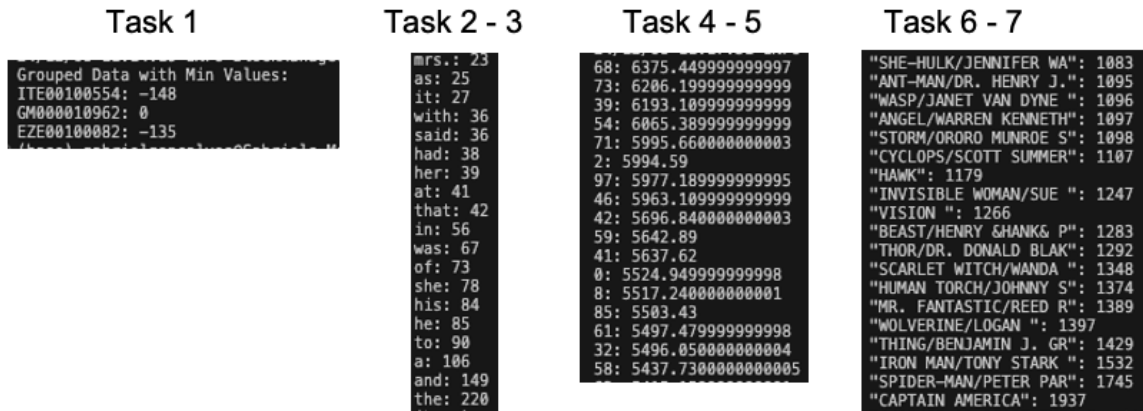


Figure 2: Outputs for all tasks in the assignment 3.

## 4 Spark SQL (Assignment 4)

### 4.1 Task 1

- Create a SparkSession.
- Define a schema for the .csv file.
- Read the CSV file into a DataFrame with the specified schema.
- Register the DataFrame as a temporary SQL table named "temps".
- Execute a SQL query to find the minimum temperature (TMIN) for each stationID.
- Collect the query results into a list of rows.
- Iterate over the results.

### 4.2 Tasks 2 and 3

- Create a SparkSession.
- Read the contents of the file Book.

- Split each line into words, remove empty strings, and create a new column word.
- Register the word\_df DataFrame as a temporary SQL table named "words".
- Execute a SQL query to count occurrences of each word, group by word, and order by count in descending order; show the results.

### 4.3 Tasks 4 and 5

- Create a Spark session.
- Read the CSV file into a DataFrame without headers.
- Assign column names to the DataFrame.
- Register the DataFrame as a temporary SQL table.
- Write an SQL query to calculate the total amount spent by each customer.
- Execute the SQL query and show results

### 4.4 Tasks 6 and 7

- Create a Spark session.
- Load the graph dataset as an RDD and count superhero occurrences.
- Convert the graph RDD to a DataFrame with column names.
- Define a schema for the names dataset.
- Load the names dataset and convert it to a DataFrame.
- Register both DataFrames as temporary SQL tables.
- Write an SQL query to join the datasets and sort by frequency.
- Execute the SQL query and show the results.

### 4.5 Terminal Outputs

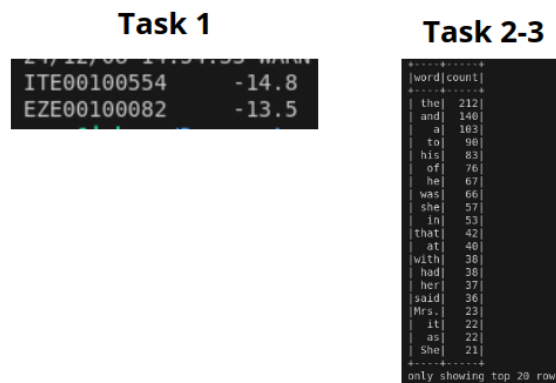


Figure 3: Outputs for tasks 1 to 3 in the assignment 4.

#### Task 4 - 5

Customer_ID	Total_Amount
68	6375.449999999997
73	6206.199999999999
39	6193.189999999999
54	6065.389999999999
71	5995.660000000003
2	5994.59
97	5977.189999999995
46	5963.189999999999
42	5696.840000000003
59	5642.89
41	5637.62
0	5524.949999999998
8	5517.240000000001
85	5503.43
61	5497.479999999998
32	5496.050000000004
58	5437.730000000005
63	5415.150000000001
15	5413.510000000001
6	5397.879999999998

only showing top 20 rows

#### Task 6 - 7

Name	Frequency
"RANDAK"	1
"MARVEL"	1
"CLUMSY"	1
"BLARE/"	1
"ZANTOR"	1
"BERSERKER"	1
"KULL"	1
"GERVASE,"	1
"JOHNSON,"	1
"SEA"	1
"DEATHCHARGE"	1
"GIURESCU,"	1
"RED"	1
"MARVEL"	1
"SHARKSKIN"	1
"FENRIS"	1
"LUNATIK"	1
"RUNE"	1
"CALLAHAN,"	1
"GOLEM"	2

only showing top 20 rows

Name	Frequency
"CAPTAIN"	1937
"SPIDER-MAN/PETER"	1745
"IRON"	1532
"THING/BENJAMIN"	1429
"WOLVERINE/LOGAN"	1397
"MR."	1389
"HUMAN"	1374
"SCARLET"	1348
"THOR/DR."	1292
"BEAST/HENRY"	1283
"VISION"	1266
"INVISIBLE"	1247
"HAWK"	1179
"CYCLOPS/SCOTT"	1107
"STORM/ORORO"	1098
"ANGEL/WARREN"	1097
"WASP/JANET"	1096
"ANT-MAN/DR."	1095
"SHE-HULK/JENNIFER"	1083
"DR."	1082

only showing top 20 rows

Figure 4: Outputs for tasks 4 to 7 in the assignment 4.